

INTRODUCTION

Here at Optimus, we spend a lot of time working with voter files, or datasets consisting of information about registered voters. (You can find more here: http://en.wikipedia.org/wiki/Voter_file) We've included a pared down voter file for you to work with (`voterfile.csv`) consisting of 50,000 records from Nevada.

PROJECT

Now that you have the data, here is the problem: Pretend it is May 2014. We need to predict voter turnout for the 2014 general election, held in November.

We would ultimately like to see four things from you:

1. Your predictions in `.csv` format (more info below).
2. A document where you interpret your results and specify the prediction methods you used.
 - a. This will be read by members of the data science team, so you can assume that we are comfortable with terminology. However, please include whatever else you think will aid us in understanding your prediction method. Your ability to communicate technical knowledge well is important as we may soon be working as part of a team!
3. A brief set of slides (no more than one cover slide and five substantive slides) where you present your results.
 - a. This should be geared toward a general audience. Use your judgment to communicate the most important points to a campaign where people likely do not have a statistics background.
4. The code that you wrote to generate your analysis.
 - a. This code should include everything that you do – reading the data in, using libraries, your analysis, any checks, renaming variables, etc.

METHOD

Use whatever prediction method you feel is appropriate, but please specify the models or algorithms that you use. Though not required, we prefer that your models or algorithms be written in a scripting language such as R, Python, or Matlab. Finally, if you manipulate the original data we give you or create new variables, please put a comment before the code that implements these changes.

.CSV RETURN FILE

The `.csv` file should contain the `optimus_id` variable plus all independent variables used to predict turnout. It should contain a field labeled `vote` with 0 if the person is predicted to stay home or 1 if they are predicted to turnout. Finally, please include a field labeled `vote_prob`, which contains the estimated probability of turning out for each individual from 0.00 to 1.00. For example it might look somewhat like the following:

optimus_id	age	vh14p	vh12g	vote	vote_prob
861681	69	1	0	1	0.729419
108469	20	0	1	1	0.635286
644435	28	0	0	0	0.238256
576830	78	0	1	1	0.604895
167371	68	1	0	1	0.752746
974034	69	0	0	0	0.164904
660415	53	1	1	1	0.768196

OVERVIEW OF VOTER FILE VARIABLES

Field	Description
optimus_id	Unique id assigned to each person
age	Age of registered voter
party	Registered political party
ethnicity	Modeled ethnicity
marital	Modeled marital status
dwellingtype	Dwelling type
education	Modeled education (commercial data)
cd	Congressional district (geography)
dma	Designated market area (geography)
occupationindustry	Modeled occupational industry (commercial data)
net_worth	Net worth (commercial data)
intrst_nascar_in_hh	Individual interested in NASCAR in household (commercial data)
petowner_dog	Likely to own a dog (commercial data)
intrst_musical_instruments_in_hh	Individual with musical interests in household (commercial data)
donates_to_liberal_causes	Donates to liberal causes (commercial data)
home_owner_or_renter	Home owner or renter (commercial data)

Voter History

Field prefixed with `vh` contain vote history information where the two digits following `vh` refer to the year of the election, and the final letter `p` or `g` indicates whether the election is a primary or general election, respectively. For example, `vh2012g` represents vote history information for the 2012 general election. A value of 1 should indicate that the individual showed up to the 2012 general election and a value of 0 should indicate that the individual did not show up to the 2012 general election.

Historical Election Returns

The remaining fields refer to historical general (`g`) or primary (`p`) precinct-level turnout numbers.

Q: How long do I have to complete this?

A: We realize you may have midterms, work, or other engagements, so we want to give you a reasonable timeframe for completing this project. **We would like for you to submit your project to us by Monday, November 9 at 11:59 p.m. to matt@optimus.com.** If you complete it earlier, you are of course welcome to submit it and we can begin reviewing early. Should anything arise and you need a bit more time, please let us know. Remember that we would like to have an answer around Thanksgiving, so keep that in mind.

Q: How long should I spend on this project?

A: That will depend entirely on you. Those of you who have experience in statistics may find the analysis easier, those of you who have experience coding may feel more comfortable working in a software package, and those of you who consider yourselves to be effective communicators may more quickly put together the write-up and presentation. Some of you may take ten hours, some of you may take a hundred hours. (Just kidding. Please don't spend 100 hours on this project. This is not meant to be that kind of project.) What I can say is that successful applicants will probably start this well before the deadline.

Q: How are we graded on this? Is it pass/fail?

A: This is not a pass/fail test. It is one of many factors we consider when selecting data science interns and fellows. We send the same exercise to both fellow and intern applicants, with obviously different expectations for both groups, but we also have your resumes and potential interviews with you. We grade each prediction exercise in terms of several facets:

- Accuracy (several different metrics for accuracy)
- Techniques/Methods
- Quality of Code
- Communication/Presentation Ability

This is not an exhaustive list - we may consider other factors as appropriate – but this should give you a good idea as to what we are looking for in a candidate. We are looking for people who think logically, solve problems creatively, and communicate effectively. How well you do is up to you!

Q: If I completely bomb this, do I still have a chance at Optimus?

A: If you're applying for a fellowship, probably not. However, if you're applying for an internship, you absolutely still have a chance if you are not able to successfully complete this task. Like you've already read, this is not pass/fail.

Q: So, what am I supposed to do again?

A: We want you to predict voter turnout for the **2014 general election**. In the voter file, this would be saved as `vh14g` (see above for explanation of voter file variables). Obviously, as it is May 2014, we do not have data for November 2014 and thus we cannot provide you with `vh14g`. We will leave it up to you as to how you choose to handle this.

Q: If I'm stuck, can I ask you questions?

A: Absolutely! The more you're able to complete on your own, the better, but we work in a team environment. Feel free to reach out to matt@optimus.com if you have questions, need clarification, or believe there to be an error.